

Lecture 2

Relational Model

Outline

- Structure of Relational Databases
- Fundamental Relational-Algebra-Operations
- Additional Relational-Algebra-Operations
- Extended Relational-Algebra-Operations

Relational Database

- Relation


instructor

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-----------|-------------|------------------|---------------|
| 10101 | Srinivasan | Comp. Sci. | 65000 |
| 12121 | Wu | Finance | 90000 |
| 15151 | Mozart | Music | 40000 |
| 22222 | Einstein | Physics | 95000 |
| 32343 | El Said | History | 60000 |
| 33456 | Gold | Physics | 87000 |
| 45565 | Katz | Comp. Sci. | 75000 |
| 58583 | Califieri | History | 62000 |
| 76543 | Singh | Finance | 80000 |
| 76766 | Crick | Biology | 72000 |
| 83821 | Brandt | Comp. Sci. | 92000 |
| 98345 | Kim | Elec. Eng. | 80000 |

Relational Database

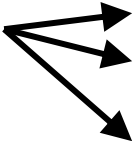
- Relation
 - Is used to refer to a table
 - conceptual representation of a table
- Attribute
 - A *column* in a *table*
- Tuple
 - A *row* in a *table*
 - A list of values

属性 **attributes**
(列 **columns**)



| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-----------|-------------|------------------|---------------|
| 22222 | Einstein | Physics | 95000 |
| 12121 | Wu | Finance | 90000 |
| 32343 | El Said | History | 60000 |
| 45565 | Katz | Comp. Sci. | 75000 |
| 98345 | Kim | Elec. Eng. | 80000 |
| 76766 | Crick | Biology | 72000 |
| 10101 | Srinivasan | Comp. Sci. | 65000 |
| 58583 | Califieri | History | 62000 |
| 83821 | Brandt | Comp. Sci. | 92000 |
| 15151 | Mozart | Music | 40000 |
| 33456 | Gold | Physics | 87000 |
| 76543 | Singh | Finance | 80000 |

元组 **tuples**
(行 **rows**)



Relational Database

- Attributes
 - Each attribute of a relation has a name
 - Domain of the attribute
 - the set of allowed values for the attribute
 - Atomicity
 - Attributes are normally required to be atomic, i.e. indivisible
 - multivalued attribute 多值属性
 - composite attribute 复合属性
 - Null
 - The special value *null* is a member of every domain

Relational Database

- Relation Instance: A specific instance of a relation (containing a specific set of rows)
- **Relation Schema**: The structure of a relation
 - *instructor* is a relation
instructor (*id*, *name*, *dept_name*, *salary*)
 - Corresponding relation schema
 $R = (id, name, dept_name, salary)$
 - *instructor* is a relation on relation schema R
instructor(R)
- In general:
 $R = (A_1, A_2, \dots, A_n)$ is a relation schema, where A_1, A_2, \dots, A_n are attributes

A relation defined over schema R is denoted by $r(R)$.

Relational Database

- Given sets of values D_1, D_2, \dots, D_n , a **relation** r is a subset of $D_1 \times D_2 \times \dots \times D_n$

| ID | name | dept_name | salary |
|-------|------------|------------|--------|
| 22222 | Einstein | Physics | 95000 |
| 12121 | Wu | Finance | 90000 |
| 32343 | El Said | History | 60000 |
| 45565 | Katz | Comp. Sci. | 75000 |
| 98345 | Kim | Elec. Eng. | 80000 |
| 76766 | Crick | Biology | 72000 |
| 10101 | Srinivasan | Comp. Sci. | 65000 |
| 58583 | Califieri | History | 62000 |
| 83821 | Brandt | Comp. Sci. | 92000 |
| 15151 | Mozart | Music | 40000 |
| 33456 | Gold | Physics | 87000 |
| 76543 | Singh | Finance | 80000 |

{22222, Einstein, Physics, 95000}

{12121, Wu, Finance, 90000}

.....

$D_1 = \{22222, 12121, \dots\}$

$D_3 = \{\text{Physics, Finance, History, Comp. Sci., } \dots\}$

$D_4 = \{95000, 90000, 85000, 80000, 75000, \dots\}$

Relational Database

- **Order of tuples is irrelevant**
 - We do not know the order of the tuples.
 - We have to specify the order if we want to visit the tuples orderly.
 - We can order the tuples by any attribute(s) that we specify.

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-----------|-------------|------------------|---------------|
| 22222 | Einstein | Physics | 95000 |
| 12121 | Wu | Finance | 90000 |
| 32343 | El Said | History | 60000 |
| 45565 | Katz | Comp. Sci. | 75000 |
| 98345 | Kim | Elec. Eng. | 80000 |
| 76766 | Crick | Biology | 72000 |
| 10101 | Srinivasan | Comp. Sci. | 65000 |
| 58583 | Califieri | History | 62000 |
| 83821 | Brandt | Comp. Sci. | 92000 |
| 15151 | Mozart | Music | 40000 |
| 33456 | Gold | Physics | 87000 |
| 76543 | Singh | Finance | 80000 |

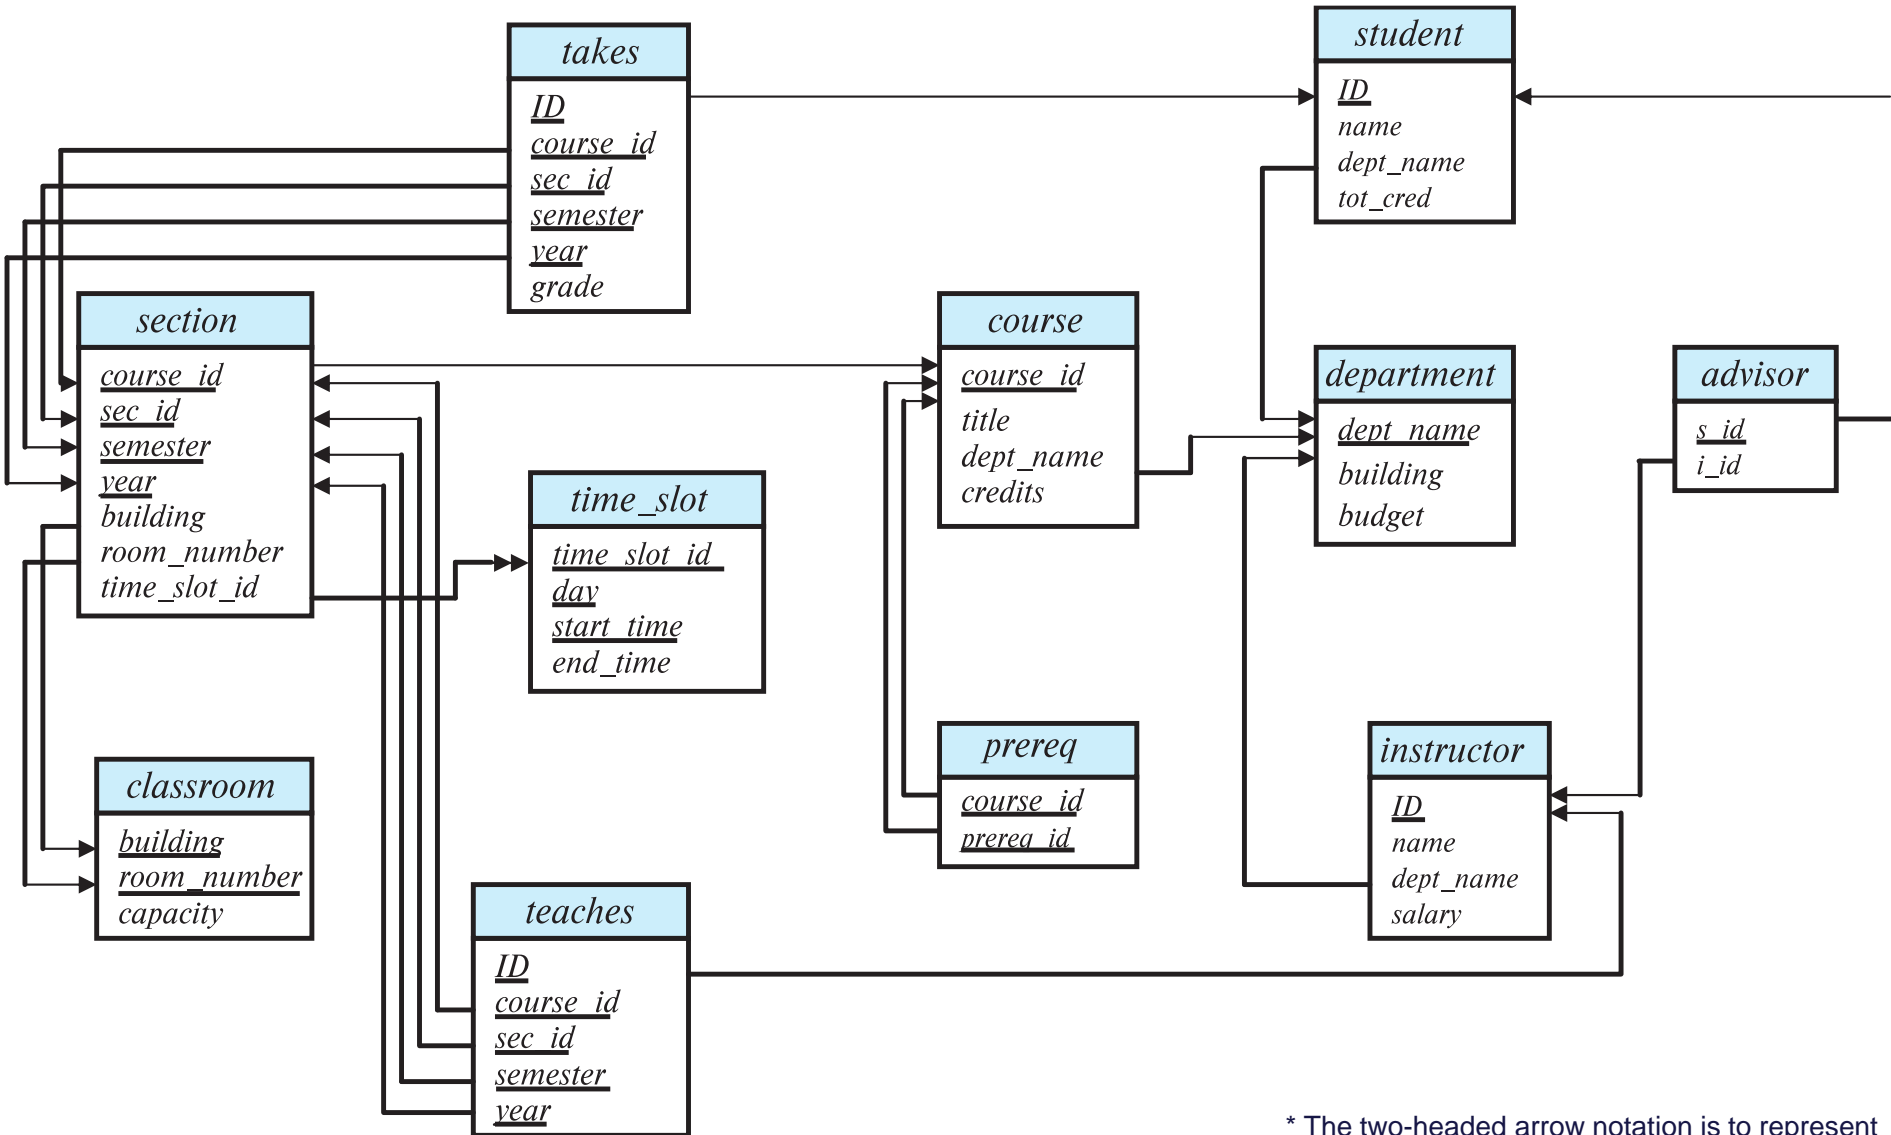
Relational Database

- Database
 - A database consists of multiple relations
 - Each relation stores one part of information
 - Information about an enterprise is broken up into parts
- student* : stores information about students**
- advisor* : stores information about which instructors
advise which students**
- instructor* : stores information about instructors**
- Why not store all information as a single relation?
 - How to break up these information reasonably?

Relational Database

- Key (键)
 - Superkey (超键)
 - Set of attributes
 - Values for which are sufficient to identify a unique tuple of each possible relation $r(R)$
 - Candidate key (候选键)
 - Is a superkey, and it is
 - Minimal (any part of which is not a superkey)
 - Primary key (主键)
 - A selected candidate key
 - Foreign key (外键)
 - Set of attributes, which are primary keys in some other relation(s)

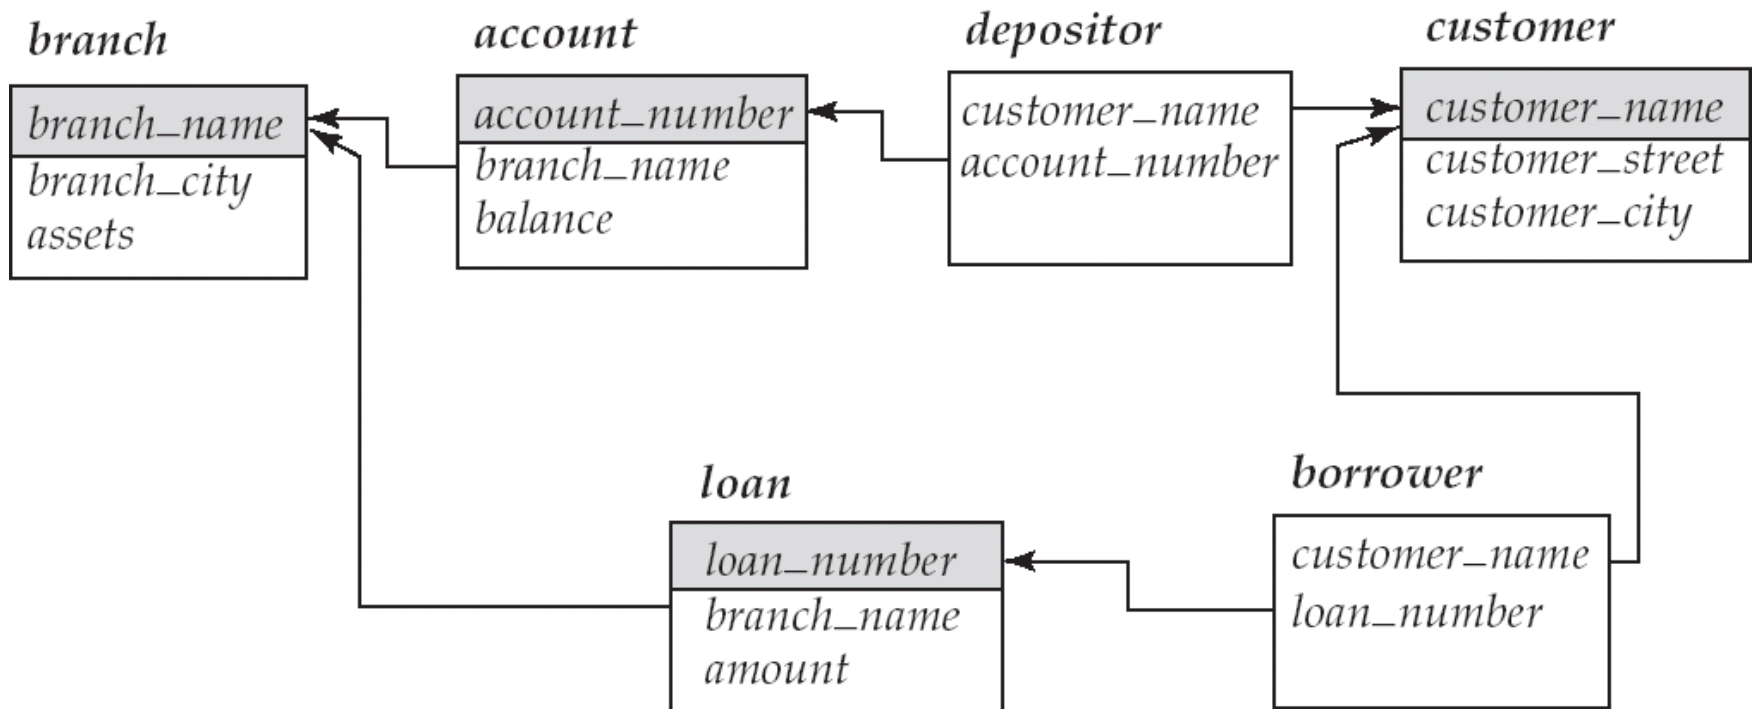
Schema Diagram



* The two-headed arrow notation is to represent a special type of referential integrity.

Schema Diagram

Another Example:



Relational Database

- Basic concepts. * SUMMARY *
 - 关系模式——属性序列
 - 关系——建立在关系模式上
 - 其值表示为集合形式，也可以表示为“表”
 - 关系实例——关系的当前值
 - 表——一般指关系的当前值，即表示关系实例
 - 属性——在关系(模式)上讨论，也可以在表上讨论
 - 在表上讨论时，可称作“列”
 - 元组——关系的（某个）元素；关系模式上并没有元组，但元组必定基于某个关系模式
 - 在表上讨论时，可称作“行”
 - (关系)数据库——一组关系，实现为一组表
 - 键

Relational Algebra

- How to make use of a database?
 - Language in which user requests information from the database
 - Query language does more than “querying” data.
- “Pure” languages:
 - Relational algebra
 - Tuple relational calculus
 - Domain relational calculus
- These three pure languages are equivalent in computing power
 - form up a basis of query languages that people use

Relational Algebra

- Procedural language
- Six basic operators
 - 选择 select: σ
 - 投影 project: Π
 - 并 union: \cup
 - 差集 set difference: $-$
 - 笛卡尔积 Cartesian product: \times
 - 重命名 rename: ρ
- Inputs: one or two relations (incl. constant relation)
- Outputs: a new relation

选择 (select)

Relation r

| A | B | C | D |
|----------|----------|----|----|
| α | α | 1 | 7 |
| α | β | 5 | 7 |
| β | β | 12 | 3 |
| β | β | 23 | 10 |

■ $\sigma_{A=B \wedge D > 5}(r)$

| A | B | C | D |
|----------|----------|----|----|
| α | α | 1 | 7 |
| β | β | 23 | 10 |

选择 (select)

- Notation: $\sigma_p(r)$
- p is called the **selection predicate** (选择谓词)
- Defined as:

$$\sigma_p(r) = \{t \mid t \in r \text{ and } p(t)\}$$

Where p is a formula in propositional calculus consisting of **terms** connected by : \wedge (**and**), \vee (**or**), \neg (**not**)

Each **term** is one of:

<attribute> op <attribute>

<attribute> op <constant>

where op is one of: $=, \neq, >, \geq, <, \leq$

- Example of selection:

$\sigma_{dept_name="Physics"}(instructor)$

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-----------|-------------|------------------|---------------|
| 22222 | Einstein | Physics | 95000 |
| 33456 | Gold | Physics | 87000 |

投影 (project)

- Relation r :

| A | B | C |
|----------|-----|-----|
| α | 10 | 1 |
| α | 20 | 1 |
| β | 30 | 1 |
| β | 40 | 2 |

■ $\Pi_{A,C}(r)$

| A | C |
|----------|-----|
| α | 1 |
| α | 1 |
| β | 1 |
| β | 2 |

=

| A | C |
|----------|-----|
| α | 1 |
| β | 1 |
| β | 2 |

投影 (project)

- Notation:

$$\Pi_{A_1, A_2, \dots, A_k}(r)$$

where A_1, A_2 are attribute names and r is a relation name.

- The result is defined as the relation of k columns obtained by erasing the columns that are not listed
- Duplicate rows removed from result, since relations are sets
- Example: To eliminate the *salary* and *id* attributes of *instructor*

$$\Pi_{name, dept_name}(instructor)$$

并 (union)

- Relations r, s :

| A | B |
|-----|-----|
|-----|-----|

| | |
|----------|---|
| α | 1 |
| α | 2 |
| β | 1 |

r

| A | B |
|-----|-----|
|-----|-----|

| | |
|----------|---|
| α | 2 |
| β | 3 |

s

$r \cup s$:

| A | B |
|-----|-----|
|-----|-----|

| | |
|----------|---|
| α | 1 |
| α | 2 |
| β | 1 |
| β | 3 |

并 (union)

- Notation: $r \cup s$
- Defined as:

$$r \cup s = \{t \mid t \in r \text{ or } t \in s\}$$

- For $r \cup s$ to be valid :
 1. r, s must have the *same* **arity** (same number of attributes)
 2. The attribute domains must be **compatible**

差集 (set difference)

- Relations r , s :

| A | B |
|----------|-----|
| α | 1 |
| α | 2 |
| β | 1 |

r

| A | B |
|----------|-----|
| α | 2 |
| β | 3 |

s

$r - s$:

| A | B |
|----------|-----|
| α | 1 |
| β | 1 |

差集 (set difference)

- Notation $r - s$
- Defined as:

$$r - s = \{t \mid t \in r \textbf{ and } t \notin s\}$$

- Set differences must be taken between **compatible relations**.
 - r and s must have the **same arity**
 - attribute domains of r and s must be **compatible**

笛卡尔积 (Cartesian product)

- Relations r , s :

| A | B |
|-----|-----|
|-----|-----|

| | |
|----------|---|
| α | 1 |
| β | 2 |

r

| C | D | E |
|-----|-----|-----|
|-----|-----|-----|

| | | |
|----------|----|-----|
| α | 10 | a |
| β | 10 | a |
| β | 20 | b |
| γ | 10 | b |

s

- $r \times s$:

| A | B | C | D | E |
|-----|-----|-----|-----|-----|
|-----|-----|-----|-----|-----|

| | | | | |
|----------|---|----------|----|-----|
| α | 1 | α | 10 | a |
| α | 1 | β | 10 | a |
| α | 1 | β | 20 | b |
| α | 1 | γ | 10 | b |
| β | 2 | α | 10 | a |
| β | 2 | β | 10 | a |
| β | 2 | β | 20 | b |
| β | 2 | γ | 10 | b |

笛卡尔积 (Cartesian product)

- Notation $r \times s$
- Defined as:

$$r \times s = \{t \ q \mid t \in r \textbf{ and } q \in s\}$$

- Assume that attributes of $r(R)$ and $s(S)$ are disjoint. (That is, $R \cap S = \emptyset$).
- If attributes of $r(R)$ and $s(S)$ are not disjoint, then **renaming** must be used.

笛卡尔积 (Cartesian product)

- Relations r, s :

| A | B |
|----------|-----|
| α | 1 |
| β | 2 |

r

| A | D | E |
|----------|-----|-----|
| α | 10 | a |
| β | 10 | a |
| β | 20 | b |
| γ | 10 | b |

s

- $r \times s$:

| $r.A$ | B | $s.A$ | D | E |
|----------|-----|----------|-----|-----|
| α | 1 | α | 10 | a |
| α | 1 | β | 10 | a |
| α | 1 | β | 20 | b |
| α | 1 | γ | 10 | b |
| β | 2 | α | 10 | a |
| β | 2 | β | 10 | a |
| β | 2 | β | 20 | b |
| β | 2 | γ | 10 | b |

Renaming

- Allows us to refer to a relation by more than one name
 - 可以通过不同的名字引用同一个关系
- Example:

$$\rho_x (E)$$

returns the expression E under the name X

- If a relational-algebra expression E has arity n , then

$$\rho_{x(A_1, A_2, \dots, A_n)} (E)$$

returns the result of expression E under the name X , and with the attributes renamed to A_1, A_2, \dots, A_n .

Renaming

- Relation r

| A | B |
|----------|-----|
| α | 1 |
| β | 2 |

r

- $r \times \rho_s(r)$

| $r.A$ | $r.B$ | $s.A$ | $s.B$ |
|----------|-------|----------|-------|
| α | 1 | α | 1 |
| α | 1 | β | 2 |
| β | 2 | α | 1 |
| β | 2 | β | 2 |

Composite Operation

- 使用多个运算构造表达式
- Example: $\sigma_{A=C}(r \times s)$
- $r \times s$

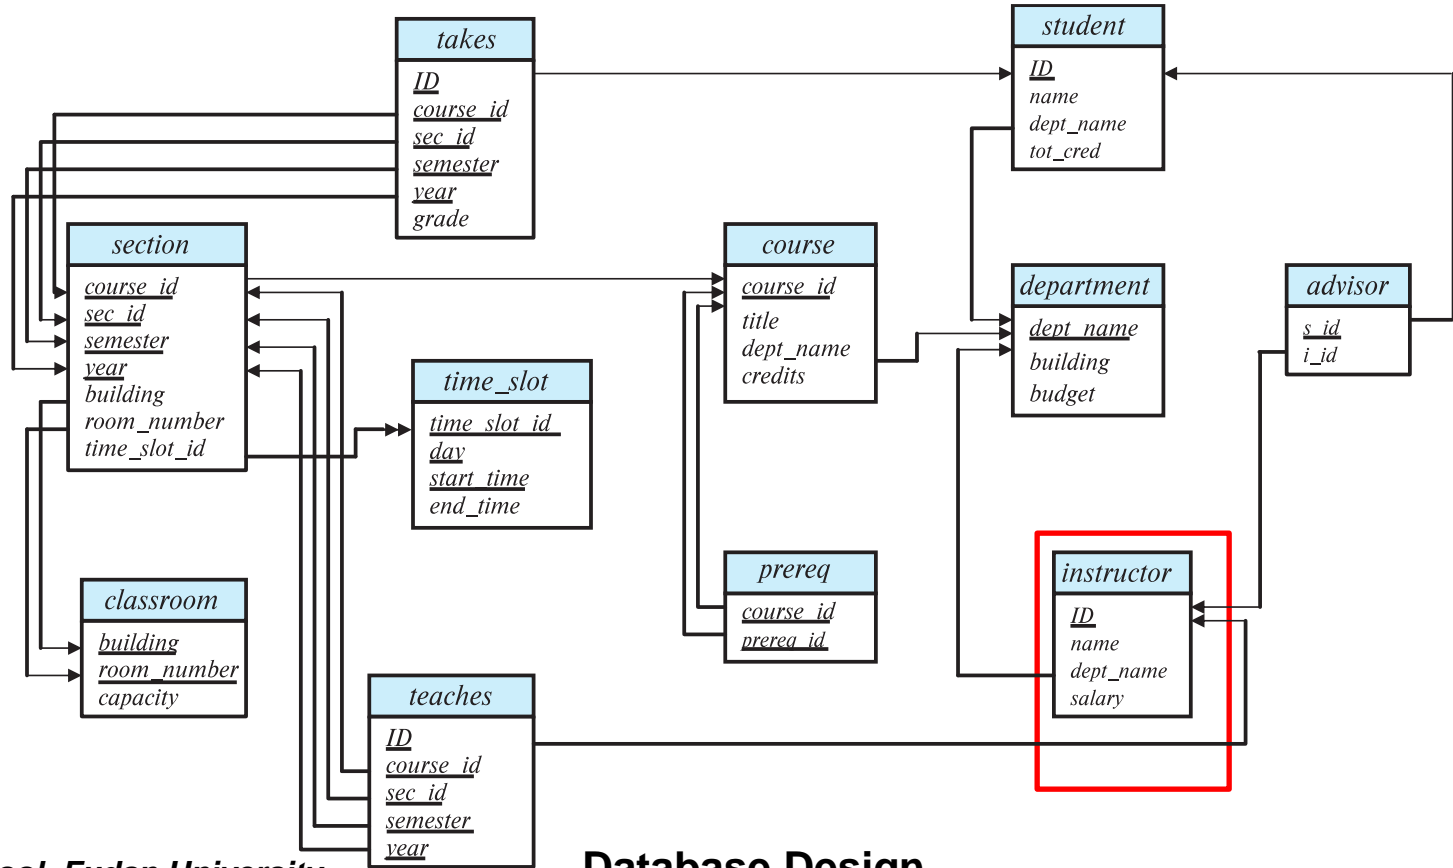
| <i>A</i> | <i>B</i> | <i>C</i> | <i>D</i> | <i>E</i> |
|----------|----------|----------|----------|----------|
| α | 1 | α | 10 | <i>a</i> |
| α | 1 | β | 10 | <i>a</i> |
| α | 1 | β | 20 | <i>b</i> |
| α | 1 | γ | 10 | <i>b</i> |
| β | 2 | α | 10 | <i>a</i> |
| β | 2 | β | 10 | <i>a</i> |
| β | 2 | β | 20 | <i>b</i> |
| β | 2 | γ | 10 | <i>b</i> |

- $\sigma_{A=C}(r \times s)$

| <i>A</i> | <i>B</i> | <i>C</i> | <i>D</i> | <i>E</i> |
|----------|----------|----------|----------|----------|
| α | 1 | α | 10 | <i>a</i> |
| β | 2 | β | 10 | <i>a</i> |
| β | 2 | β | 20 | <i>b</i> |

Example Queries - *University*

- Select and Project
 - Find instructors with salary greater than \$85,000
 - Find the ID and salary of the instructors with salary greater than \$85,000

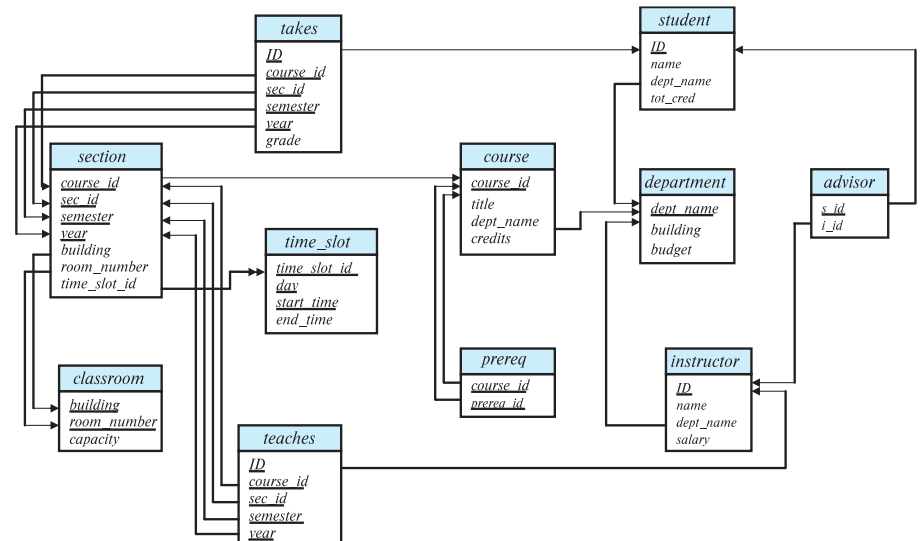


Example Queries - *University*

- Select and Project
 - Find instructors with salary greater than \$85,000
 - Find the ID and salary of the instructors with salary greater than \$85,000

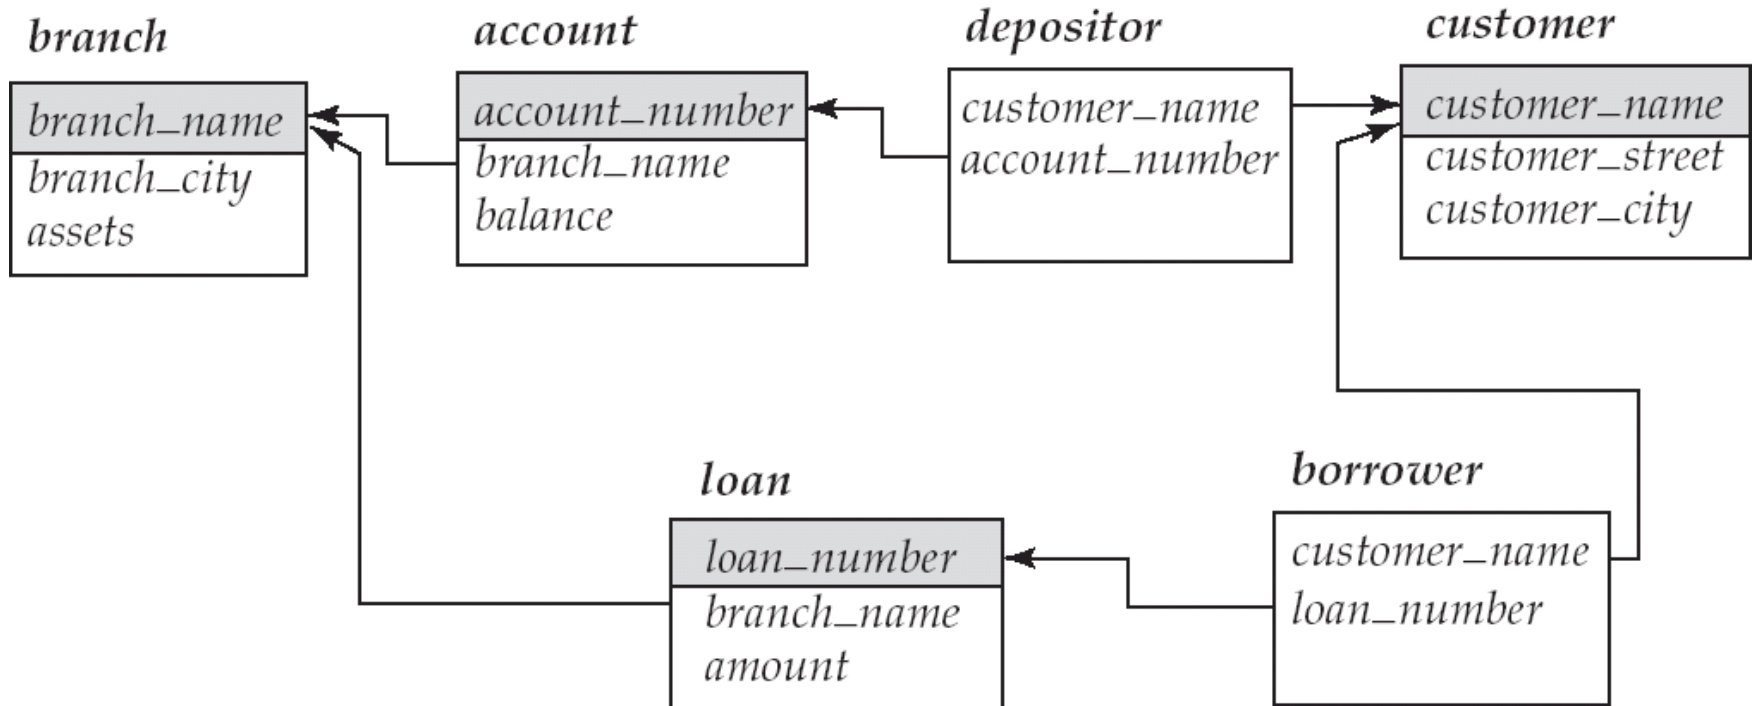
$\sigma_{salary > 85000} (instructor)$

$\Pi_{id, salary} (\sigma_{salary > 85000} (instructor))$



Banking Example

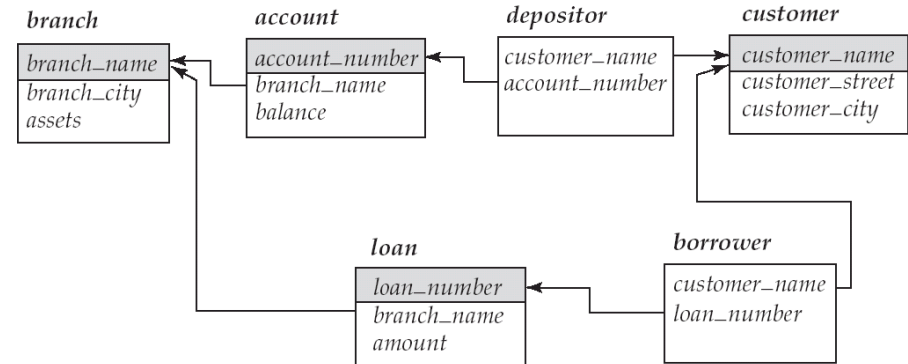
Schema diagram:



Example Queries

- Union and Set-Difference

- 拥有贷款或者存款的客户名
- 拥有存款但没有贷款的客户名
- 拥有存款和贷款的客户名



$$\Pi_{customer_name}(borrower) \cup \Pi_{customer_name}(depositor)$$

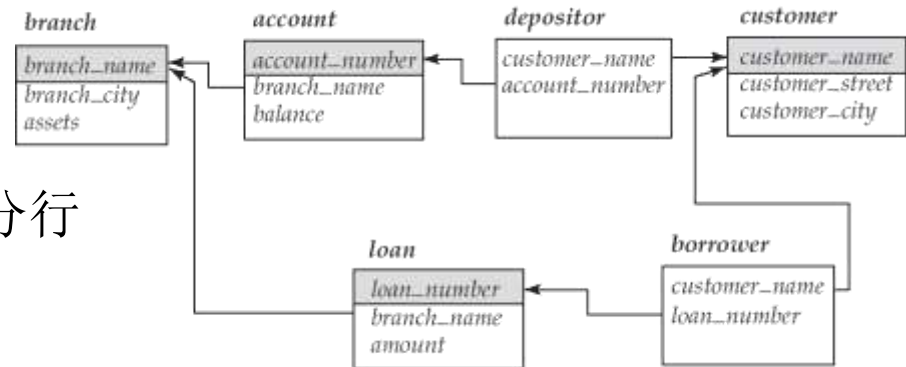
$$\Pi_{customer_name}(depositor) - \Pi_{customer_name}(borrower)$$

$$\Pi_{customer_name}(depositor) - (\Pi_{customer_name}(depositor) - \Pi_{customer_name}(borrower))$$

Example Queries

- Cartesian-Product

- 在Perryridge分行有贷款的客户名
- 在Perryridge分行有贷款且在任何分行都没有存款的客户名



$$\Pi_{customer_name} (\sigma_{branch_name="Perryridge"} (\sigma_{borrower.loan_number = loan.loan_number} (borrower \times loan)))$$

$$\Pi_{customer_name} (\sigma_{branch_name = "Perryridge"}$$

$$(\sigma_{borrower.loan_number = loan.loan_number} (borrower \times loan))) -$$

$$\Pi_{customer_name} (depositor)$$

Example Queries

- 关系运算的组合顺序
 - 在Perryridge分行有贷款的客户名
 - 先算borrower x loan

$\Pi_{customer_name} (\sigma_{branch_name = "Perryridge"} (\sigma_{borrower.loan_number = loan.loan_number} (borrower \times loan)))$

- 先对loan进行branch_name="Perryridge"选择

$\Pi_{customer_name} (\sigma_{loan.loan_number = borrower.loan_number} (\sigma_{branch_name = "Perryridge"} (loan) \times borrower))$

Basic Operators - Summary

- Let E_1 and E_2 be relational-algebra expressions; the following are all relational-algebra expressions:
 - $E_1 \cup E_2$
 - $E_1 - E_2$
 - $E_1 \times E_2$
 - $\sigma_p(E_1)$, P is a predicate on attributes in E_1
 - $\Pi_S(E_1)$, S is a list consisting of some of the attributes in E_1
 - $\rho_x(E_1)$, x is the new name for the result of E_1

Additional Operations

- 交集 (Intersection)
- 自然联结 (Natural-join)
- 除 (Division)
- 赋值 (Assignment)

交集 (intersection)

- Relation r, s :

| A | B |
|----------|---|
| α | 1 |
| α | 2 |
| β | 1 |

r

| A | B |
|----------|---|
| α | 2 |
| β | 3 |

s

- $r \cap s$

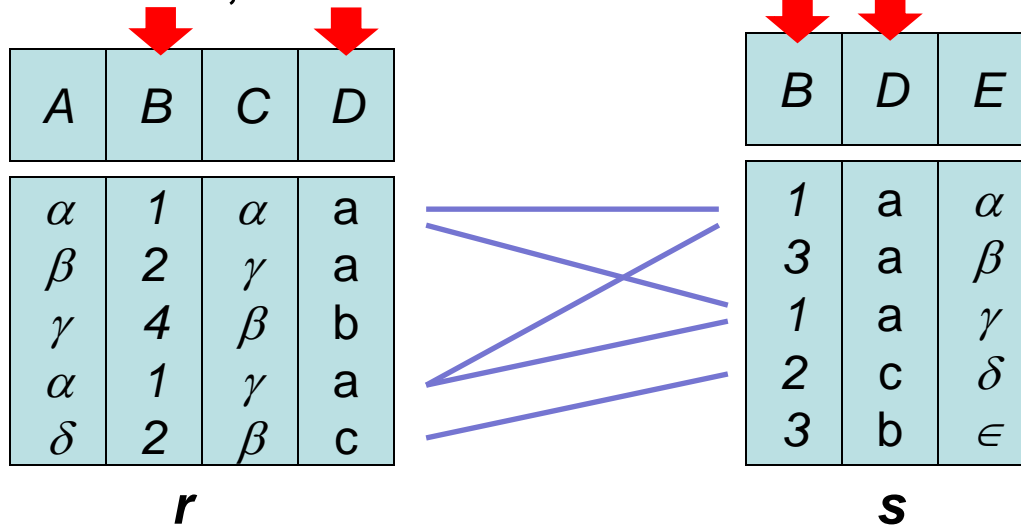
| A | B |
|----------|---|
| α | 2 |

交集 (intersection)

- Notation: $r \cap s$
- Defined as:
- $r \cap s = \{ t \mid t \in r \textbf{ and } t \in s \}$
- Assume:
 - r, s have the *same arity*
 - attributes of r and s are compatible
- Note: $r \cap s = r - (r - s)$

自然联结 (natural-join)

- Relations r, s :



$r \bowtie s$

| A | B | C | D | E |
|----------|-----|----------|-----|----------|
| α | 1 | α | a | α |
| α | 1 | α | a | γ |
| α | 1 | γ | a | α |
| α | 1 | γ | a | γ |
| δ | 2 | β | c | δ |

自然联结 (natural-join)

- Let r and s be relations on schemas R and S respectively.
Then, natural-join $r \bowtie s$ is a relation on schema $R \cup S$ (“union”) obtained as follows:
 - Consider each pair of tuples t_r from r and t_s from s .
 - If t_r and t_s have the same value on each of the attributes in $R \cap S$ (“intersection”), add a tuple t to the result, where
 - t has the same value as t_r on r
 - t has the same value as t_s on s

- Example:

$R = (A, B, C, D)$

$S = (E, B, D)$

- Result schema = (A, B, C, D, E)
- $r \bowtie s$ is defined as:

$$\Pi_{r.A, r.B, r.C, r.D, s.E} (\sigma_{r.B = s.B \wedge r.D = s.D} (r \times s))$$

Theta Join

- The **join** operation allows us to combine a ***select*** operation and a ***Cartesian-Product*** operation into a single operation.
- The Theta Join operation is a generalized form of natural join.
- Consider relations $r (R)$ and $s (S)$
- Let “theta” be a predicate on attributes in the schema R “union” S. The join operation $r \bowtie_{\theta} s$ is defined as follows:

$$r \bowtie_{\theta} s = \sigma_{\theta} (r \times s)$$

- Thus

$$\sigma_{instructor.id = teaches.id} (instructor \times teaches)$$

Can equivalently be written as

$$instructor \bowtie_{instructor.id = teaches.id} teaches$$

- In this case, above query is equivalent to $instructor \bowtie teaches$

除 (division)

Relations r, s :

| A | B |
|-----|-----|
|-----|-----|

| | |
|---------------|---|
| α | 1 |
| α | 2 |
| α | 3 |
| β | 1 |
| γ | 1 |
| δ | 1 |
| δ | 3 |
| δ | 4 |
| ε | 6 |
| ε | 1 |
| β | 2 |

r

| B |
|-----|
|-----|

| |
|---|
| 1 |
| 2 |

s

$r \div s$:

| A |
|-----|
|-----|

| |
|----------|
| α |
| β |

除 (division)

- Let r and s be relations on schemas R and S respectively where

$$R = (A_1, \dots, A_m, B_1, \dots, B_n)$$

$$S = (B_1, \dots, B_n)$$

The result of $r \div s$ is a relation on schema

$$R - S = (A_1, \dots, A_m)$$

$$r \div s = \{ t \mid t \in \Pi_{R-S}(r) \wedge \forall u \in s (tu \in r) \}$$

Where tu means the concatenation of tuples t and u to produce a single tuple

除 (division)

Relations r , s :

| A | B | C | D | E |
|-----|-----|-----|-----|-----|
|-----|-----|-----|-----|-----|

| | | | | |
|----------|-----|----------|-----|-----|
| α | a | α | a | 1 |
| α | a | γ | a | 1 |
| α | a | γ | b | 1 |
| β | a | γ | a | 1 |
| β | a | γ | b | 3 |
| γ | a | γ | a | 1 |
| γ | a | γ | b | 1 |
| γ | a | β | b | 1 |

r

| D | E |
|-----|-----|
|-----|-----|

| | |
|-----|-----|
| a | 1 |
| b | 1 |

s

$r \div s$:

| A | B | C |
|-----|-----|-----|
|-----|-----|-----|

| | | |
|----------|-----|----------|
| α | a | γ |
| γ | a | γ |

赋值 (assignment)

- Notation \leftarrow
 - 使用临时变量保存关系运算的结果
 - 简化关系表达式的表示
- $r \div s$ may be expressed as:

$$temp1 \leftarrow \Pi_{R-S}(r)$$

$$temp2 \leftarrow \Pi_{R-S}((temp1 \times s) - \Pi_{R-S,S}(r))$$

$$r \div s = temp1 - temp2$$

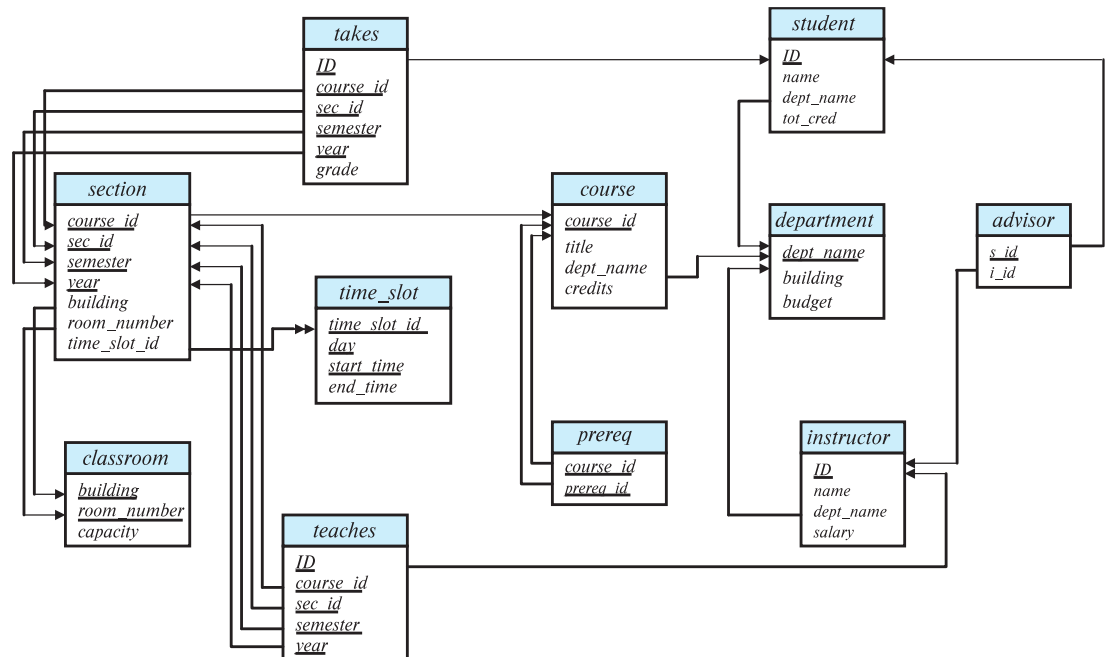
- The result to the right of the \leftarrow is assigned to the relation variable on the left of the \leftarrow
- May use variable in subsequent expressions

Example Queries

- Set Intersection
 - Find the ids of the courses taught in Fall 2009 as well as in Spring 2010

$\Pi_{course_id} (\sigma_{semester = 'Fall' \wedge year = 2009} (section))$

$\cap \Pi_{course_id} (\sigma_{semester = 'Spring' \wedge year = 2010} (section))$



Example Queries

- Natural-join
 - Find the names of instructors and the corresponding ids of the courses taught by the instructors

$$\Pi_{name, course_id}(instructor \bowtie teaches)$$

- Find the names of instructors and the corresponding ids of the courses taught by the instructors, and show the year in which the courses were taught

$$\Pi_{name, course_id, year}(instructor \bowtie teaches)$$

Example Queries

- Division
 - Find the names of instructors and the corresponding ids of the courses taught by the instructors in both 2010 and 2011

$$\Pi_{name, course_id, year} (instructor \bowtie teaches) \\ \div \rho_{temp(year)} (\{(2010), (2011)\})$$

- 一般的查询方法

$$\Pi_{name, course_id} (\sigma_{year = 2010} (instructor \bowtie teaches)) \\ \cap \Pi_{name, course_id} (\sigma_{year = 2011} (instructor \bowtie teaches))$$

- What does the following query mean?

$$\Pi_{name, course_id, year} (instructor \bowtie teaches) \\ \div \Pi_{year} (\sigma_{semester = \text{“Fall”}} (section))$$

Review

- 基本的关系代数运算
 - 六种基本运算符，其中：
 - 一元运算符和二元运算符各三种
- 附加的关系代数运算
 - 能够由基本关系代数运算表达
 - 丰富和简化了关系代数的表达方式

Extended Operations

- 泛化投影 (Generalized Projection)
- 聚合函数 (Aggregate Functions)
- 外联结 (Outer Join)

Generalized Projection

$$\Pi_{F_1, F_2, \dots, F_n}(E)$$

E is any relational-algebra expression

Each of F_1, F_2, \dots, F_n are arithmetic expressions involving constants and attributes in the schema of E

- Extended project operation
 - allowing arithmetic functions to be used in the projection list
- Example
 - Relation ***credit_info(customer_name, limit, credit_balance)***
Find how much more each person can spend :

$$\Pi_{customer_name, limit - credit_balance}(credit_info)$$

聚合 (Aggregation)

- **Aggregation function (聚合函数)**: takes a collection of values and returns a single value as a result

avg: average value

min: minimum value

max: maximum value

sum: sum of values

count: number of values

- **Aggregate operation** in relational algebra

$$G_1, G_2, \dots, G_n \mathcal{g} F_1(A_1), F_2(A_2), \dots, F_n(A_n) (E)$$

E is any relational-algebra expression

- G_1, G_2, \dots, G_n is a list of attributes on which to group (can be empty)
- Each F_i is an aggregate function
- Each A_i is an attribute name

Also use the symbol γ : $G_1, G_2, \dots, G_n \gamma F_1(A_1), F_2(A_2), \dots, F_n(A_n) (E)$

Aggregation Examples

- Relation r :

| A | B | C |
|----------|----------|-----|
| α | α | 7 |
| α | β | 7 |
| β | β | 3 |
| β | β | 10 |

$g_{\text{sum}(C)}(r)$

| $\text{sum}(C)$ |
|-----------------|
| 27 |

Aggregation Examples

- Relation *instructor* grouped by *dept_name*:

| <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-------------|------------------|---------------|
| Einstein | Physics | 95000 |
| Gold | Physics | 87000 |
| Wu | Finance | 90000 |
| Singh | Finance | 80000 |
| Mozart | Music | 40000 |

dept_name Σ *sum(salary)* (*instructor*)

| <i>dept_name</i> | sum(salary) |
|------------------|--------------------|
| Physics | 182000 |
| Finance | 170000 |
| Music | 40000 |

聚合 (Aggregation)

- Renaming the result of aggregation
 - Result of aggregation does not have a name
 - For convenience, we permit renaming as part of aggregate operation

dept_name **g** sum(*salary*) **as sum_salary** (*instructor*)

| <i>dept_name</i> | <i>sum_salary</i> |
|------------------|-------------------|
| Physics | 182000 |
| Finance | 170000 |
| Music | 40000 |

Outer Join

- Relation *course*

| <i>course_id</i> | <i>title</i> | <i>dept_name</i> | <i>credits</i> |
|------------------|--------------|------------------|----------------|
| BIO-301 | Genetics | Biology | 4 |
| CS-190 | Game Design | Comp. Sci. | 4 |
| CS-315 | Robotics | Comp. Sci. | 3 |

- Relation *prereq*

| <i>course_id</i> | <i>prereq_id</i> |
|------------------|------------------|
| BIO-301 | BIO-101 |
| CS-190 | CS-101 |
| CS-347 | CS-101 |

- Inner Join (Natural Join)
course ⋈ *prereq*

| <i>course_id</i> | <i>title</i> | <i>dept_name</i> | <i>credits</i> | <i>prereq_id</i> |
|------------------|--------------|------------------|----------------|------------------|
| BIO-301 | Genetics | Biology | 4 | BIO-101 |
| CS-190 | Game Design | Comp.Sci. | 4 | CS-101 |

Outer Join

- Relation *course*

| <i>course_id</i> | <i>title</i> | <i>dept_name</i> | <i>credits</i> |
|------------------|--------------|------------------|----------------|
| BIO-301 | Genetics | Biology | 4 |
| CS-190 | Game Design | Comp. Sci. | 4 |
| CS-315 | Robotics | Comp. Sci. | 3 |

- Relation *prereq*

| <i>course_id</i> | <i>prereq_id</i> |
|------------------|------------------|
| BIO-301 | BIO-101 |
| CS-190 | CS-101 |
| CS-347 | CS-101 |

- Left Outer Join
course \bowtie *prereq*

| <i>course_id</i> | <i>title</i> | <i>dept_name</i> | <i>credits</i> | <i>prere_id</i> |
|------------------|--------------|------------------|----------------|-----------------|
| BIO-301 | Genetics | Biology | 4 | BIO-101 |
| CS-190 | Game Design | Comp. Sci. | 4 | CS-101 |
| CS-315 | Robotics | Comp. Sci. | 3 | null |

Outer Join

- Relation *course*

| <i>course_id</i> | <i>title</i> | <i>dept_name</i> | <i>credits</i> |
|------------------|--------------|------------------|----------------|
| BIO-301 | Genetics | Biology | 4 |
| CS-190 | Game Design | Comp. Sci. | 4 |
| CS-315 | Robotics | Comp. Sci. | 3 |

- Relation *prereq*

| <i>course_id</i> | <i>prereq_id</i> |
|------------------|------------------|
| BIO-301 | BIO-101 |
| CS-190 | CS-101 |
| CS-347 | CS-101 |

- Right Outer Join
course \bowtie_{right} *prereq*

| <i>course_id</i> | <i>title</i> | <i>dept_name</i> | <i>credits</i> | <i>prere_id</i> |
|------------------|--------------|------------------|----------------|-----------------|
| BIO-301 | Genetics | Biology | 4 | BIO-101 |
| CS-190 | Game Design | Comp. Sci. | 4 | CS-101 |
| CS-347 | <i>null</i> | <i>null</i> | <i>null</i> | CS-101 |

Outer Join

- Relation *course*

| <i>course_id</i> | <i>title</i> | <i>dept_name</i> | <i>credits</i> |
|------------------|--------------|------------------|----------------|
| BIO-301 | Genetics | Biology | 4 |
| CS-190 | Game Design | Comp. Sci. | 4 |
| CS-315 | Robotics | Comp. Sci. | 3 |

- Relation *prereq*

| <i>course_id</i> | <i>prereq_id</i> |
|------------------|------------------|
| BIO-301 | BIO-101 |
| CS-190 | CS-101 |
| CS-347 | CS-101 |

- Full Outer Join
course \bowtie *prereq*

| <i>course_id</i> | <i>title</i> | <i>dept_name</i> | <i>credits</i> | <i>prere_id</i> |
|------------------|--------------|------------------|----------------|-----------------|
| BIO-301 | Genetics | Biology | 4 | BIO-101 |
| CS-190 | Game Design | Comp. Sci. | 4 | CS-101 |
| CS-315 | Robotics | Comp. Sci. | 3 | <i>null</i> |
| CS-347 | <i>null</i> | <i>null</i> | <i>null</i> | CS-101 |

Null Values (空值)

- Denoted by *null*
- *null* signifies an unknown value or that a value does not exist.
- The result of any arithmetic expression involving *null* is *null*.
- Aggregate functions simply ignore null values (as in SQL)
- For duplicate elimination and grouping, null is treated like any other value, and two nulls are assumed to be the same (as in SQL)

Null Values (空值)

- Three-valued logic using the truth value *unknown*:
 - OR: (*unknown* **or** *true*) = *true*,
(*unknown* **or** *false*) = *unknown*
(*unknown* **or** *unknown*) = *unknown*
 - AND: (*true* **and** *unknown*) = *unknown*,
(*false* **and** *unknown*) = *false*,
(*unknown* **and** *unknown*) = *unknown*
 - NOT: (**not** *unknown*) = *unknown*
- Result of select predicate is treated as *false* if it evaluates to *unknown*

Modification of the Database

- The content of the database may be modified using the following operations
 - 添加(insertion)
 - 删除(deletion)
 - 修改(updating)
- Expressed using the assignment operator

添加 (insertion)

$$r \leftarrow r \cup E$$

where r is a relation and E is a relational algebra expression

- To insert data into a relation, we either:
 - specify a tuple to be inserted (E is a constant relation, containing only one tuple)
 - write a query whose result is a set of tuples to be inserted
- 例

$$instructor \leftarrow instructor \cup \{(54321, \text{"Mickey"}, \text{Music}, 45000)\}$$

删除 (deletion)

$$r \leftarrow r - E$$

where r is a relation and E is a relational algebra query

- Selected tuples are removed from the database.
 - similarly to a query
 - can delete only whole tuples
 - cannot delete values on only particular attributes
- 例

$$instructor \leftarrow instructor - \sigma_{dept_name = "Music"}(instructor)$$

$$student \leftarrow student - \sigma_{tot_cred \geq 1 \wedge tot_cred \leq 5}(student)$$

删除 (deletion)

- 例

- 删除位于Needham市的分行的所有存款帐户

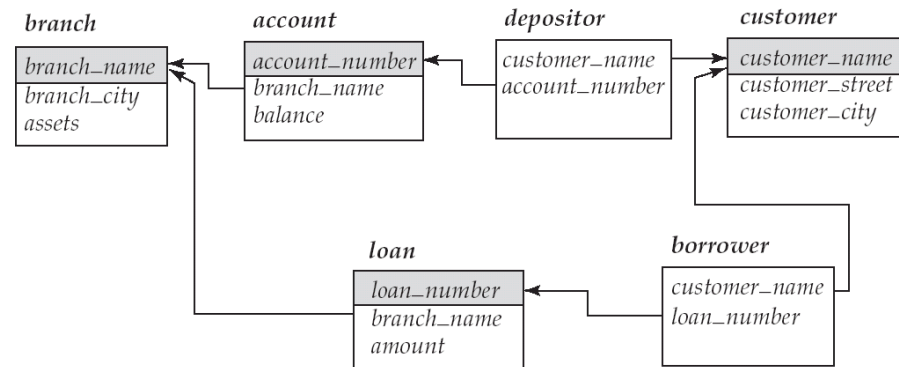
$r_1 \leftarrow \sigma_{branch_city = "Needham"} (account \bowtie branch)$

$r_2 \leftarrow \Pi_{account_number, branch_name, balance} (r_1)$

$r_3 \leftarrow \Pi_{customer_name, account_number} (r_2 \bowtie depositor)$

$account \leftarrow account - r_2$

$depositor \leftarrow depositor - r_3$



更新 (Updating)

$$r \leftarrow \prod_{F_1, F_2, \dots, F_l} (r)$$

其中 r 是关系,

F_i 是

- r 的第 i 个属性 (如果第 i 个属性不需要修改); 或
- 仅包含常量和 r 中属性的算术表达式, 根据此值更新 r 的第 i 个属性

Each F_i is either

- **the i^{th} attribute of r , if the i^{th} attribute is not updated, or,**
 - **if the attribute is to be updated, an expression, involving only constants and the attributes of r , which gives the new value for the attribute**
- Change a value in a tuple without changing *all* values in the tuple

更新 (Updating)

- 例
 - 向存款帐户发放利息，金额为存款帐户余额的5%

$account \leftarrow \Pi_{account_number, branch_name, balance * 1.05} (account)$

- 对余额超过10000元的存款帐户支付6%的红利，其他存款帐户支付5%的红利

$account \leftarrow$

$\Pi_{account_number, branch_name, balance * 1.06} (\sigma_{balance > 10000} (account))$
 $\cup \Pi_{account_number, branch_name, balance * 1.05} (\sigma_{balance \leq 10000} (account))$

小结

- 什么是关系数据库？
- 关系代数的六个基本操作以及其他的关系运算
- 采用关系代数表达式表示数据库查询和修改操作
 - 简单查询
 - 组合查询
 - 增、删、改
- 外联结在数据库查询中的作用

Next Lecture

- SQL
 - 从纯语言到应用
 - 在实践中学习各类查询语言的用法

End of Lecture 2