

## Lecture 4

# Intermediate SQL

# Outline

- Join Expressions
- Views
- Transactions
- SQL Data Types and Schemas
- Integrity Constraints
- Authorization

# Joined Relations

- **Join operations** take two relations and return as a result another relation.
- A join operation is a **Cartesian product** which requires that **tuples in the two relations match** (under some condition).
  - It also specifies the attributes that are present in the result of the join
- The join operations are typically used as subquery expressions in the **from** clause

*Given relations:  $r(A1, A2, A3)$   $s(A1, A4, A5)$*

**select  $r.A1, A4$  from  $r$  natural join  $s$**

# Join operations – Example

- Relation *course*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

- Relation *prereq*

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

- Observe that  
prereq information is missing for CS-315 and  
course information is missing for CS-347

# Outer Join

- An extension of the join operation that **avoids loss of information**.
- Computes the join and then **adds tuples** from one relation **that does not match** tuples in the other relation to the result of the join.
- Uses *null* values.

# Left Outer Join

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

*course* **natural left outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>

# Right Outer Join

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

*course* **natural right outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

# Full Outer Join

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

*course* **natural full outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101



# Joined Relations

- **Join operations** take two relations and return as a result another relation.
- These additional operations are typically used as subquery expressions in the **from** clause
- **Join condition** – defines which tuples in the two relations match, and what attributes are present in the result of the join.
- **Join type** – defines how tuples in each relation that do not match any tuple in the other relation (based on the join condition) are treated.

<i>Join types</i>	<i>Join Conditions</i>
<b>inner join</b> <b>left outer join</b> <b>right outer join</b> <b>full outer join</b>	<b>natural</b> <b>on</b> <predicate> <b>using</b> ( $A_1, A_2, \dots, A_n$ )

# Joined Relations – Examples

- **course inner join prereq on**  
*course.course\_id = prereq.course\_id*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>	<i>course_id</i>
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190

- course left outer join prereq on**  
*course.course\_id = prereq.course\_id*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>	<i>course_id</i>
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190
CS-315	Robotics	Comp. Sci.	3	<i>null</i>	<i>null</i>

# Joined Relations – Examples

*course* **natural right outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

*course* **full outer join** *prereq* **using** (*course\_id*)

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

# Joined Relations - Examples

- List all departments along with the number of instructors in each department

```
select department.dept_name, count(distinct instructor.id)
from department left outer join instructor on
department.dept_name=instructor.dept_name
group by department.dept_name;
```

Oracle:

```
select department.dept_name, count(distinct instructor.id)
from department, instructor
where department.dept_name = instructor.dept_name(+)
group by department.dept_name;
```

# Views

- In some cases, it is not desirable for all users to see the entire logical model (that is, all the actual relations stored in the database.)
- Consider a person who needs to know an instructors name and department, but not the salary. This person should see a relation described, in SQL, by

```
select ID, name, dept_name  
from instructor
```

- A **view** provides a mechanism to hide certain data from the view of certain users.
- Any relation that is not of the conceptual model but is made visible to a user as a “virtual relation” is called a **view**.

# View Definition

- A view is defined using the **create view** statement which has the form

create view **v** as < query expression >

where <query expression> is any legal SQL expression.  
The view name is represented by *v*.

- Once a view is defined, the view name can be used to refer to the virtual relation that the view generates.
- View definition **is not** the same as creating a new relation by evaluating the query expression
  - Rather, a view definition causes the **saving of an expression**; the expression is *substituted* into queries using the view.

# Example Views

- A view of instructors without their salary  
**create view** *faculty* **as**  
    **select** *ID, name, dept\_name*  
    **from** *instructor*
- Find the names of all instructors in the Biology department  
**select** *name*  
**from** *faculty*  
**where** *dept\_name* = 'Biology'
- Create a view of department salary totals  
**create view** *departments\_total\_salary*(*dept\_name*,  
*total\_salary*) **as**  
    **select** *dept\_name, sum (salary)*  
    **from** *instructor*  
    **group by** *dept\_name*;

# Views Defined Using Other Views

- One view may be used in the expression defining another view
  - A view relation  $v_1$  is said to *depend directly* on a view relation  $v_2$  if  $v_2$  is used in the expression defining  $v_1$
  - A view relation  $v_1$  is said to *depend on* view relation  $v_2$  if either  $v_1$  depends directly to  $v_2$  or there is a path of dependencies from  $v_1$  to  $v_2$
  - A view relation  $v$  is said to be *recursive* if it depends on itself.



# Views Defined Using Other Views

- **create view** *physics\_fall\_2009* **as**  
    **select** *course.course\_id, sec\_id, building, room\_number*  
    **from** *course, section*  
    **where** *course.course\_id = section.course\_id*  
          **and** *course.dept\_name = 'Physics'*  
          **and** *section.semester = 'Fall'*  
          **and** *section.year = '2009';*
- **create view** *physics\_fall\_2009\_watson* **as**  
    **select** *course\_id, room\_number*  
    **from** *physics\_fall\_2009*  
    **where** *building= 'Watson';*

# View Expansion

- A way to define the meaning of views defined in terms of other views.
  - View expansion of an expression repeats the following replacement step:  
**repeat**  
    Find any view relation  $v_i$  in  $e_1$   
    Replace the view relation  $v_i$  by the expression defining  $v_i$   
**until** no more view relations are present in  $e_1$ 
    - As long as the view definitions are not recursive, this loop will terminate

# View Expansion

- Expand use of a view in a query/another view

```
create view physics_fall_2009_watson as  
(select course_id, room_number  
from (select course.course_id, sec_id, building, room_number  
        from course, section  
        where course.course_id = section.course_id  
            and course.dept_name = 'Physics'  
            and section.semester = 'Fall'  
            and section.year = '2009')  
where building = 'Watson';
```



*physics\_fall\_2009*

# Update of a View

- Create a view of instructors without their salary  
**create view *faculty* as**  
**select *ID, name, dept\_name***  
**from *instructor***
- Add a new tuple to the *faculty* view  
**insert into *faculty* values ('30765', 'Green', 'Music');**

This insertion must be represented by the insertion of the tuple

**('30765', 'Green', 'Music', null)**

into the *instructor* relation.

# Some Updates cannot be Translated Uniquely

- **create view** *instructor\_info* **as**  
    **select** *ID, name, building*  
    **from** *instructor, department*  
    **where** *instructor.dept\_name= department.dept\_name;*
- **insert into** *instructor\_info* **values** ('69987', 'White', 'Taylor');
  - which department, if multiple departments in Taylor?
  - what if no department is in Taylor?
- Most SQL implementations allow updates only on simple views
  - The **from** clause has only one database relation.
  - The **select** clause contains only attribute names of the relation, and does not have any expressions, aggregates, or **distinct** specification.
  - Any attribute not listed in the **select** clause can be set to null
  - The query does not have a **group** by or **having** clause.

# And Some Not at All

- **create view** *history\_instructors* **as**  
    **select** \*  
    **from** *instructor*  
    **where** *dept\_name*= 'History';
- What happens if we insert ('25566', 'Brown', 'Biology', 100000) into *history\_instructors*?

# Materialized Views \*\*

- **Materializing a view**: create a physical table containing all the tuples in the result of the query defining the view
- If relations used in the query are updated, the materialized view result becomes out of date
  - Need to **maintain** the view, by updating the view whenever the underlying relations are updated.

# Transactions

- Unit of work
- Atomic transaction
  - either fully executed or rolled back as if it never occurred
- Isolation from concurrent transactions
- Transactions begin implicitly
  - Ended by **commit work** or **rollback work**
- But default on most databases: each SQL statement commits automatically
  - Can turn off auto commit for a session (e.g., using an API)
  - In SQL:1999, can use: **begin atomic .... end**



# Built-in Data Types

- Basic Data Types
  - char(n), varchar(n)
  - int , smallint, numeric(p,d)
  - real, double precision, float(n)
- Addition Built-in Data Types
  - date: **date** '2006-09-04'
  - time: **time** '08:55:00'
  - timestamp: **timestamp** '2006-09-04 11:25:09.77'
  - interval: Indicates a period of time, usu. implemented as a fixed point number
    - Subtracting a date/time/timestamp value from another gives an interval value
    - Interval values can be added to date/time/timestamp values

# Built-in Data Types

- Operating with date/time/timestamp values
  - Extract values of individual fields  
`extract (year from r.starttime)`
  - Type coercion  
`cast '2006-12-25' as date`
- Implementation may vary in real systems
  - Oracle supports addition/subtraction between datetime types and numeric values

# User-Defined Types

- Creating types or domains
  - Examples:
    - **create type** Dollars **as** numeric(12,2) final
    - **create domain** person\_name **as** char(20) not null
    - Then we can specify attributes with types and domains (if they are supported by DBMS)
  - Domains can have constraints specified on them while types cannot
  - Domain constraints are the most elementary form of integrity constraint

# User-Defined Types

- Domain Constraints
  - New domains can be created from existing data types
  - Domains are not strongly typed, while types are.
    - Example

```
create type USD numeric(12, 2)
create type CNY numeric(12, 2)
```

We cannot assign or compare a value of type USD to a value of type CNY.
- Type coercion is applicable.  
(**cast** r.A **as** CNY)

# Large-Object Types

- Large objects (photos, videos, etc.) are stored as a *large object*
  - **blob**: binary large object – stored as uninterpreted binary data
  - **clob**: character large object – stored as characters
  - When a query returns a large object, a *pointer* is returned rather than the large object itself.
    - Returned data should be handled by an application **outside of** the database system

# Integrity Constraints

- Integrity constraints **guard against accidental damage** to the database, by ensuring that authorized changes to the database do not result in a loss of data consistency.
  - A checking account must have a balance ***greater*** than \$10,000.00
  - A salary of a bank employee must be ***at least*** \$4.00 an hour
  - A customer must ***have*** a (*non-null*) phone number

# Integrity Constraints

- Constraints on a Single Relation
  - **not null**
  - **primary key**
  - **unique**
  - **check( $P$ )**
- Referential Integrity
  - **foreign key**
- Assertion across relations

# Not Null and Unique Constraints

- **not null**
  - Declare *name* and *budget* to be **not null**  
*name* **varchar**(20) **not null**  
*budget* **numeric**(12,2) **not null**
- **unique** (  $A_1, A_2, \dots, A_m$  )
  - The unique specification states that the attributes  $A_1, A_2, \dots, A_m$  form a candidate key.
  - Candidate keys are permitted to be null (in contrast to primary keys).



# The **check** clause

- **check** (P)


where P is a predicate

Example: ensure that semester is one of fall, winter, spring or summer:

```
create table section (  
    course_id varchar (8),  
    sec_id varchar (8),  
    semester varchar (6),  
    year numeric (4,0),  
    building varchar (15),  
    room_number varchar (7),  
    time slot id varchar (4),  
    primary key (course_id, sec_id, semester, year),  
    check (semester in ('Fall', 'Winter', 'Spring', 'Summer'))  
);
```


# Referential Integrity

- Ensures that a value that appears in one relation for a given set of attributes also appears for a certain set of attributes in another relation.
  - Example: If “Biology” is a department name appearing in one of the tuples in the *instructor* relation, then there exists a tuple in the *department* relation for “Biology” department
- Let  $A$  be a set of attributes. Let  $R$  and  $S$  be two relations that contain attributes  $A$  and where  $A$  is the primary key of  $S$ .  $A$  is said to be a **foreign key** of  $R$  if for any values of  $A$  appearing in  $R$  these values also appear in  $S$ .



# Cascading Actions in Referential Integrity

- **create table** *course* (  
    ...  
    *dept\_name* **varchar**(20),  
    **foreign key** (*dept\_name*) **references** *department*  
        **on delete cascade**  
        **on update cascade**,  
    ...  
)



- alternative actions to **cascade**: **set null**, **set default**
  - e.g., on delete set null

# Assertions

- An **assertion** is a predicate expressing a condition that we wish the database always to satisfy.
- An assertion in SQL takes the form  
**create assertion** <assertion-name> **check** <predicate>
- The system tests an assertion
  - when an assertion is made
  - when an update may violate the assertion
  - that may introduce a significant amount of overhead
    - USE with GREAT CARE!

Some implementation of DBMS does not support assertion.  
Refer to the user's manual for instructions.

# Authorization

- Authorizations to ...
  - read data
  - insert new data
  - update data
  - delete data
- Privileges
  - select
  - insert
  - update
  - delete

# Authorization Specification in SQL

- The **grant** statement
  - grant** <privilege list>  
**on** <relation name or view name> **to** <user / role list>
  - several privileges can be granted in one command
  - <privilege list> may be **all privileges**, indicating all allowable privileges to be granted
  - <user / role list> may be **public**, allowing all current and future users the privilege granted *implicitly*
  - Granting a privilege on a view does not imply granting any privileges on the underlying relations.
  - The grantor of the privilege must already hold the privilege on the specified item (or be the database administrator).

# Revoking Authorization in SQL

- The **revoke** statement
  - revoke** <privilege list>  
**on** <relation name or view name> **from** <user list>
  - <privilege list> may be all to revoke all privileges revokee may hold
  - <revokee-list> may include **public**, meaning all users lose the privilege except those granted it *explicitly*
  - If the same privilege was granted twice to the same user by different grantees, the user may retain the privilege after the revocation

# Privileges in SQL

- **select:** allows read access to relation, or the ability to query using the view
  - Example: grant users  $U_1$ ,  $U_2$ , and  $U_3$  **select** authorization on the *branch* relation:

**grant select on *branch* to  $U_1, U_2, U_3$**

- **insert:** the ability to insert tuples
- **update:** the ability to update using the SQL update statement
- **delete:** the ability to delete tuples.
- **references:** ability to declare foreign keys when creating relations.
- **all privileges:** used as a short form for all the allowable privileges



# Privilege To Grant Privileges

- **with grant option**: allows a user who is granted a privilege to pass the privilege on to other users.
  - Example:  
**grant select on *branch* to  $U_1$  with grant option**  
gives  $U_1$  the **select** privileges on *branch* and allows  $U_1$  to grant this privilege to others

# Roles

- Roles permit common privileges for a class of users can be specified just once by creating a corresponding “role”
- Privileges can be granted to or revoked from roles, just like user
- Roles can be assigned to users, and even to other roles
  - **create role** instructor;
  - **grant** *instructor* **to** Amit;
  - **grant select on** *takes* **to** *instructor*;
  - **create role** *teaching\_assistant*
  - **grant** *teaching\_assistant* **to** *instructor*;
    - *Instructor* inherits all privileges of *teaching\_assistant*
- Chain of roles

# Authorization and Views

- Users can be given authorization on **views**, *without* being given any authorization on the ***relations*** used in the view definition
- Ability of views to hide data serves both
  - to **simplify usage** of the system and
  - to **enhance security** by allowing users access only to data they need for their job
- A combination of relational-level security and view-level security can be used to precisely limit a user's access to the data that user needs.

# Authorization on Views

- **create view** *geo\_instructor* **as**  
  (**select** \*  
  **from** *instructor*  
  **where** *dept\_name* = 'Geology');
- **grant select on** *geo\_instructor* **to** *geo\_staff*
- Suppose that a *geo\_staff* member issues
  - **select** \*  
  **from** *geo\_instructor*;
- What if
  - *geo\_staff* does not have permissions on *instructor*?
  - creator of view did not have some permissions on *instructor*?

# Authorization on Views


- Creation of view does not require **resources** authorization since no real relation is being created
- The creator of a view gets only those privileges that provide no additional authorization beyond that he already had.
  - e.g. if creator of view *geo\_instructor* had only **read** authorization on *instructor*, he gets only **read** authorization on *geo\_instructor*

# Other Authorization Features

- **references** privilege to create foreign key
  - **grant reference** (*dept\_name*) **on department** **to** Mariano;
  - why is this required?
- transfer of privileges
  - **grant select on department** **to** Amit **with grant option**;
  - **revoke select on department from** Amit, Satoshi **cascade**;
  - **revoke select on department from** Amit, Satoshi **restrict**;


# Granting of Privileges

- The passage of authorization from one user to another may be represented by an authorization graph.
- The nodes of this graph are the users.
- The root of the graph is the database administrator.
- Graph below
  - An edge  $U_i \rightarrow U_j$  indicates that user  $U_i$  has granted update authorization on loan to  $U_j$ .




# Authorization Grant Graph

- *Requirement:* All edges in an authorization graph must be part of some path originating with the database administrator
- If DBA revokes grant from  $U_1$ :
  - Grant must be revoked from  $U_4$  since  $U_1$  no longer has authorization
  - Grant must not be revoked from  $U_5$  since  $U_5$  has another authorization path from DBA through  $U_2$





- Must prevent cycles of grants with no path from the root:
  - DBA grants authorization to  $U_7$
  - $U_7$  grants authorization to  $U_8$
  - $U_8$  grants authorization to  $U_7$
  - DBA revokes authorization from  $U_7$
- Must revoke grant  $U_7$  to  $U_8$  and from  $U_8$  to  $U_7$  since there is no path from DBA to  $U_7$  or to  $U_8$  anymore.



# Limitations of SQL Authorization

- SQL does not support authorization ***at a tuple level***
  - E.g. we cannot restrict students to see only (the tuples storing) **their own** grades
- With the growth in Web access to databases, database accesses come primarily from *application servers*.
  - End users don't have database user ids; they are all mapped to **the same database user id**
  - All end-users of an application (such as a web application) may be mapped to a single database user

# Limitations of SQL Authorization

- The task of authorization in above cases falls on the application program, with no support from SQL
  - Benefit: fine grained authorizations, such as to individual tuples, can be implemented by the application.
  - Drawback: Authorization must be done in application code, and may be dispersed all over an application
  - Checking for absence of authorization loopholes becomes very difficult since it requires reading large amounts of application code

# Next Lecture

- Advanced SQL

End of Lecture 4